

### 1. Why do we need chunking in NLP and RAG?

Chunking is needed because LLMs have a token limit, and searching long documents efficiently requires breaking them into smaller parts. Without chunking:

- The model might miss key information.
  - Retrieval could bring back irrelevant or incomplete context.
  - Large documents may exceed token limits, making them hard to process.
- 

### 2. Why use a vector database instead of SQL for AI-powered search?

A vector database (like Pinecone, Weaviate, or FAISS) is designed for **semantic search**, meaning it retrieves **similar** data based on meaning, not exact matches.

#### Key Differences:

Feature	SQL	Vector DB
Data Type	Structured (tables, rows)	Unstructured (embeddings)
Search Type	Exact match	Similarity search
Speed	Slow for high-dimensional data	Fast for high-dimensional data
Use Case	Traditional queries	AI-powered retrieval

---

### 3. What is prompt engineering, and why is it important?

Prompt engineering is designing effective instructions for LLMs to get the best results. It's important because **how** you ask influences **what** you get.

#### Examples of Techniques:

- **Zero-shot prompting** – Directly ask the model a question.
  - **Few-shot prompting** – Give examples before the query.
  - **Chain-of-thought prompting** – Encourage step-by-step reasoning.
  - **Role-based prompting** – Ask the model to act as an expert.
- 

### 4. What are inference parameters in LLMs, and why use temperature = 1?

Inference parameters control text generation behavior.

#### Key Parameters:

- **Temperature** (0–1+) → Controls randomness.
  - **Low (0-0.5)**: More predictable.
  - **High (0.8-1.5)**: More creative.
  - **1.0 (Balanced)**: Mix of creativity and coherence.
- **Top-k sampling** → Picks from the top K likely words.
- **Top-p (nucleus sampling)** → Picks from the top p% of probable words.

Using **temperature = 1** gives a mix of **creativity** and **coherence** in responses.

---

## 5. How to develop a Multimodal RAG system for text & images?

A **Multimodal RAG** retrieves both text and image data.

**Architecture:**

1. **Text embeddings** → Using models like **OpenAI's text-embedding-ada-002** or **BERT**.
2. **Image embeddings** → Using models like **CLIP** or **BLIP**.
3. **Fusion model** → Combines text + image data.
4. **Vector database** → Stores both embeddings.

**Tech Stack:**

- **LangChain** for RAG pipeline.
  - **FAISS/Pinecone** for storage.
  - **CLIP** for image retrieval.
- 

## 6. Is OCR the best approach for extracting text in a multimodal RAG system?

**Limitations of OCR:**

- Struggles with **handwritten text** and **poor-quality images**.
- **Misses layout context** (tables, figures).
- **No semantic understanding** (only extracts text).

**Alternatives:**

- **LayoutLM** → Extracts text + layout info.
- **Donut** → Directly generates structured output without OCR.

- **Tesseract OCR + LLM** → Hybrid approach for better accuracy.
- 

## 7. How to generate embeddings from transactional data efficiently?

### Steps:

1. Convert the table into meaningful **text-based representations**.
  - Example: "User A bought Item X for \$50 on Jan 1"
2. Use **transformers (like BERT, T5, or Sentence-BERT)** to create embeddings.
3. Store embeddings in a **vector database**.

### Scaling for millions of rows:

- Use **FAISS/HNSW indexing** for fast retrieval.
  - Process data in **batches** instead of all at once.
  - Use **distributed computing (Spark, Ray, Dask)** for parallel processing.
- 

## 8. How to summarize a 100-page PDF efficiently?

### Approach:

1. **Split the PDF into chunks** (e.g., by paragraphs).
2. **Use an embedding model** (e.g., OpenAI's text embeddings) to find key sections.
3. **Summarize each chunk** using models like **T5, GPT-4, or BART**.
4. **Combine & refine** summaries into a final version.

### Tools:

- **LangChain** (for chunking & retrieval).
  - **Hugging Face models** (for summarization).
  - **LLMs (GPT-4, Claude)** for final refinements.
- 

## 9. Difference between GPT-3.5 and GPT-4?

### Main Differences:

Feature	GPT-3.5	GPT-4
---------	---------	-------

Accuracy	Good	Better
Reasoning	Decent	Stronger
Multimodal	No	Yes (handles text + images)
Context Window	Shorter	Longer

### **Why GPT-4 is better?**

- **More accurate answers.**
- **Improved logic and reasoning.**
- **Can process images (GPT-4V).**